

A Sublinear Space, Polynomial Time Algorithm for Directed s-t Connectivity

Greg Barnes, Jonathan F. Buss, Walter L. Ruzzo, Baruch Schieber


STCON in Directed Unique-Path Graphs

Sampath Kannan, Sanjeev Khanna, Sudeepa Roy

STCON = $\{(G, s, t) | t \text{ is reachable from } s \text{ in directed graph } G\}$

USTCON = $\{(G, s, t) | s, t \text{ are connected in undirected graph } G\}$

- USTCON \in L [Reingold,05]
- STCON \in DSPACE($\log^2 n$) [Savitch's theorem]
- BFS/DFS solves STCON using $O(n \log n)$ space and $O(m+n)$ time.
- STCON can be solved using $\frac{n}{2^{\Theta(\sqrt{\log n})}}$ space and polynomial time. [BBRS,98]



D_i := Set of vertices in i -th level of BFS tree

Fix an integer λ .

Divide the levels D_0, D_1, \dots, D_{n-1} of the bfs tree rooted at s into λ equivalence classes $C_0, C_1, \dots, C_{\lambda-1}$ where $C_i = \{v \in V(G) \mid \text{dist}(s, v) = i \pmod{\lambda}\}$.

There exists some C_i with number of vertices $\leq \frac{n}{\lambda}$

Lemma

Every vertex reachable from s is at distance at most λ from $C_i \cup \{s\} \forall i \in \{0, 1, \dots, \lambda - 1\}$.

Proof.

Let v be a vertex reachable from s such that $\text{dist}(s, v) = \lambda q + r$

where $0 \leq q \leq \frac{n}{\lambda}, 0 \leq r \leq \lambda - 1$.

Consider the shortest path P from s to v .

If $r \geq i$, then let u be the vertex on P such that $\text{dist}(s, u) = \lambda q + i$. Otherwise let u be the vertex on P such that $\text{dist}(s, u) = \max\{0, \lambda(q - 1) + i\}$. Then $u \in C_i$ and $\text{dist}(u, v) \leq \lambda$. ■

Algorithm BFS(G,s,t):

for $i=0$ to $\lambda - 1$:

$C_i = \{s\}$

for every vertex v at distance i from s :

if $|C_i| < \frac{n}{\lambda}$:

Add v to C_i

else:

$i := i+1$

for $q=1$ to $\frac{n}{\lambda}$:

$S := \varnothing$

for every vertex v at distance λ from C_i :

if $|C_i| + |S| < \frac{n}{\lambda}$:


Add v to S

else:

$i := i+1$

$C_i := C_i \cup S$

Check if t is within distance λ from C_i



Space complexity: $O(\frac{n \log n}{\lambda} + S_{PATH}(\lambda, n))$



Time complexity: $O(\frac{n^3}{\lambda} \cdot T_{PATH}(\lambda, n))$

$f(n)$ -bounded STCON: Given an n -vertex directed graph G and two vertices s, t of G , check whether t is within distance $f(n)$ of s

Idea:

Fix an integer k . We divide the vertices into k sets according to their vertex number modulo k .

Suppose we want to find all paths from s of length $\leq L = f(n)$. We can enumerate all $(L+1)$ -digit k -ary numbers $(s \bmod k, a_1, a_2, \dots, a_L)$ and for each such number, find all L -length paths $(s, u_1, u_2, \dots, u_L)$ such that $(s, u_1, u_2, \dots, u_L) \bmod k = (s \bmod k, a_1, a_2, \dots, a_L)$.



Algorithm for L-bounded STCON(G,s,t,L):

Set $V_0 := \{s\}$

for all (L+1)-digit k-ary number $(s \bmod k, a_1, a_2, \dots, a_L)$:

for $i=1$ to L:

for all $u \in V_{i-1}$ AND $v = a_i \bmod k$:

if (u,v) is an edge:

Add v to V_i



V_i contains all vertices reachable from s using a path that maps to $(s \bmod k, a_1, a_2, \dots, a_L)$ modulo k .

if $v=t$:

Return "Connected!"

Erase V_{i-1} from work tape

We don't need to maintain the list of vertices V_i . Instead we can maintain a $\left\lceil \frac{n}{k} \right\rceil$ -bit vector where j -th entry of the vector will be equal to 1 if $(kj + a_i) \in V_i$. Then the algorithm will run in $O(\frac{n}{k} + L \log k)$ space and $O(k^L L \cdot \frac{n^2}{k^2})$ time.



While computing V_i , instead of looking for vertices at distance 1 from V_{i-1} , we can look for vertices at distance L^{r-1} recursively and that would give us an algorithm for L^r -bounded STCON.

Short paths recursive algorithm:

SPR($G, L, r, s \bmod k, t \bmod k, V_s$): V_s is a subset of vertices with vertex number = $s \bmod k$

if $r=0$:

for all $u \in V_s$ and $v = t \bmod k$:

if (u,v) is an edge:

add v to V_t

else:

$V_0 := V_s$

for all $(L+1)$ -digit k -ary number $(s \bmod k, a_1, a_2, \dots, a_L)$:

for $i=1$ to L :

$V_i = \text{SPR}(G, L, r-1, a_{i-1} \bmod k, a_i \bmod k, V_{i-1})$

Erase V_{i-1} from work tape

Return $V_t \cap V_L$

Then $S_{PATH}(n, L^r) = O(\frac{n}{k} + L \log k) + S_{PATH}(n, L^{r-1}) \implies S_{PATH}(n, L^r) = O(r(\frac{n}{k} + L \log k))$

and $T_{PATH}(n, L^r) = O(k^L L) \cdot T_{PATH}(n, L^{r-1})$ where $T(n, 1) = \frac{n^2}{k^2}$

$$\implies T_{PATH}(n, L^r) = O((k^L L)^r \cdot \frac{n^2}{k^2})$$

If we take $\lambda = L^r$, then space complexity of the first algorithm is

$S(n) = O(\frac{n \log n}{L^r} + r(\frac{n}{k} + L \log k))$ and time complexity is

$$T(n) = O(\frac{n^3}{L^r} \cdot k^{Lr} L^r \cdot \frac{n^2}{k^2}) = O(n^5 k^{Lr-2})$$

Taking $L=4$, $r = \sqrt{\log n}$, $k = 2^{\Theta(\sqrt{\log n})}$ gives us $S(n) = \frac{n}{2^{\Theta(\sqrt{\log n})}}$ and $T(n) = n^{O(1)}$.



- STCON in regular digraphs can be solved using $O(\log n)$ space. [RTV,06]
- STCON in planar DAGs with single source can be solved using $O(\log n)$ space. [ABCDR,06]
- STCON in directed unique-path graphs is solvable using $O(n^\epsilon \log n)$ space and $O(n^{\frac{1}{\epsilon}})$ time. [KKR,08]

Definition

A directed graph G is a unique-path graph w.r.t. a source vertex s if there is at most one simple path from s to every vertex $v \in V(G)$. For every vertex x , $N^+(x) = \{v \in V(G) \mid (x,v) \in E(G)\}$ and $N^-(x) = \{v \in V(G) \mid (v,x) \in E(G)\}$.

Remark



DFS-tree of a unique-path graph doesn't have any forward edges or cross edges.



$\forall v \in V(G), next_v : N(v) \rightarrow N(v)$ is a bijective function.

Edge successor function is defined as $succ(u,v) = (v, next_v(u))$.

Then the sequence $e, succ(e), succ^2(e), succ^3(e), \dots$ is a dfs traversal of the tree.



Algorithm:

Step 1:

Current vertex x is either a newly discovered vertex or DFS has backtracked to x from some $v \in N^+(x)$

Initially $x = s$

if x is a newly discovered vertex:

 Ask the oracle for the first vertex in $N^+(x)$

else if DFS has backtracked from v to x :

 Ask the oracle for successor of v in $N^+(x)$

Step 2:

if oracle says there are no more vertices in $N^+(x)$:

 Perform the **backtrack step** find $\text{Parent}[x]$, set current vertex $x := \text{Parent}[x]$ and go back to step 1

else if oracle returns a vertex $y \in N^+(x)$:

 Perform the **discovery step** to check whether (x,y) is a back-edge.

 if (x,y) is a back edge: \\ This is the only bad case

 then ask the oracle for successor of y in $N^+(x)$ and go back to step 2

 else:

 if $y == t$:

 Return "t is reachable from s"

 else:

 Set current vertex $x := y$ and go back to step 1

For a unique-path graph, if we store the entire active path, we can perform the backtrack step and the discovery step easily.

L-bounded DFS is a DFS which backtracks whenever length of the active path exceeds L.

Lemma

In a unique-path graph, L-bounded DFS can be implemented in $O(L \log n)$ space and $O(n+mL)$ time.

$O(\sqrt{n} \log n)$ -space algorithm for DFS on unique-path graphs:

Maintaining landmark vertices: In stead of storing the entire active path, we will store landmark vertices which are a few evenly spaced vertices on the active path from s to current vertex x. The landmark vertices will be $s = z_0, z_1, \dots, z_{p-1}, z_p = x$ where for $0 \leq i \leq p-1$, z_i is the vertex at distance $i \lfloor \sqrt{n} \rfloor$ from s on the current active path.

Backtrack step from current vertex x :

for all $v_i \in N^-(x)$:



Perform a \sqrt{n} -bounded DFS
from z_{p-1} in $G \setminus (v_i, x)$

If x isn't reached in this DFS:

Return v_i as Parent[x]

Space: $O(\sqrt{n} \log n)$

Time: $O(mn + m^2\sqrt{n})$



Discovery step from current vertex x:

Oracle has returned a vertex y and we have to check whether (x,y) is a back edge i.e. whether y has been visited before.


Let $Z(y)$ be the set of landmark vertices reachable from y in \sqrt{n} -bounded DFS.

If y has been visited before, then there exist two landmark vertices z_{j-1} and z_j such that y lies between these two landmark vertices on the active path.

Lemma

$z_k \notin Z(y)$ for $k > j$ i.e. z_j is the landmark vertex of highest index in $Z(y)$.

Proof.

Otherwise there are two distinct simple paths from y to z_k - one given by the \sqrt{n} -bounded DFS and other one is along the active path via z_j . 

Discovery step for edge (x,y) :

if $y \in \{z_0, z_1, \dots, z_{p-1}\}$:

Return "y has been visited before"

Perform \sqrt{n} -bounded DFS from y

if $Z(y) == \varnothing$:

Return "y has not been visited before"

Let j be the highest index such that $z_j \in Z(y)$

if $j == 0$:

Return "y has not been visited before"

if $j \leq p - 1$:

Perform \sqrt{n} -bounded DFS from z_{j-1} and terminate as soon as z_j or y is reached

if y is reached:

Return "y has been visited before"

else:

Return "y has not been visited before"

if $j == p$:

Perform $2\sqrt{n}$ -bounded DFS from z_{p-2} and terminate as soon as x or y is reached

if y is reached:

Return "y has been visited before"

else:

Return "y has not been visited before"

Space: $O(\sqrt{n} \log n)$

Time: $O(mn + m^2\sqrt{n})$

Hence we get a $O(\sqrt{n} \log n)$ space and polynomial time algorithm for STCON in unique-path graphs.

This can be improved to $O(n^\epsilon)$ space and $n^{O(\frac{1}{\epsilon})}$ time algorithm for all $\epsilon > 0$.